

LLM latency by backend × K (bars: ms; annotated 'not run' rows)

